


[Polls](#) [Participants](#) [Reports](#) [Teams](#) [Upgrades](#) [Help](#)
[Presenter tips](#)
[bobkillen881](#)

# KubeCon / Cloud Native Con NA 2019 RUG

## Survey Results

Current run (last updated Nov 23, 2019 11:40am)

5

Polls


21

Participants

26

Average responses

### Which industry or group do you represent?



Response options

Count Percentage

**Academia****9 43%****21**

Responses

Government

4 19%

Non Profit

1 5%

Private Sector

4 19%

Other

3 14%

### Are you using Kubernetes to support research workloads?



Response options

Count Percentage

Not currently.

2 13%

**15**

Responses

Not currently, but looking into it.


3 20%

Yes, in development or testing purposes.

4 27%

**Yes, in production.****6 40%**

# What tools or frameworks are you using or would like to use?



## Responses

Keycloak   Kubernetes   cilogin   GitLab   CI/CD   singularity

Singularity   slurm   OSG   Opa   opa   spark   Spark

Kubeflow   pytorch   tensorflow

Kubernetes, Dask, Spark, HTCondor, KubeFlow, Jupyter, Binder   Opa

jupyter   Kind   Jupyterhub   jupyter jupyterlab tensorflow

htcondor   kubeflow   Jupyter

Multi-cloud/multi-k8s (on-prem, hosted) infrastructure for supporting national laboratory user and facility web services

spark   Jupyter   Kubernetes, kube-batch

Kubeadm, jupyter, gitlab, Prometheus, kata, containerd, helm, helmfile, oauth2-proxy, mlflow, nginx-ingress, grafana, elasticsearch

Would like: Service mesh once we can wrap our head around it.

Data as a service

Oh, I guess we are doing some of this at work. We've got a GPU cluster set up and we're trying to enable it as a GPUaaS cluster.

Cloud events, JupyterHub, data workflows, data analysis hpc job scheduling

openshift   OpenShift   K8s

Rancher (for multitenancy & integrated tools (Grafana, etc)). OPA to apply security policies.


No clue. Completely new to this space. I'm here for me, not for work, so I'm mostly just checking it out. But I'm very interested in research and could see a career change in that direction.

Using Kubespray, interested in Volcano   Home Grown   K8s   Kubeflow

43

Responses

# What is the biggest pain point you've encountered with using Kubernetes as a research platform?




Responses

How to apply policy	Almost too many options & new added every quarter		
Multi Cluster	Learning Curve for our Users	It's easy*	
run-to-completion job support	IPV6	Learning Curve	
"effort required"	Scheduling	Shared file systems	Networking
Team knowledge	accounting	culture-change	
file-system integration!	Security	Learning curve	Learningcurve
Scheduling	Understanding and translating customer's needs		slurm
OrganizationalSupport	security	munge-key	multi-tenancy
Learning Curve	multi-cluster	scaling	lack of scheduling features
Multitenancy	networking	scheduling	Multi-tenancy
			scale

35

Responses

# Where should we focus our efforts?



Responses

Documentation
Usability for End Users
batch
multitenancy
Shared filesystems
multi-cluster
End user docs
Software lifecycles (upgrading software)
Multi-tenancy
push enhancements upstream
Security
Batch
batch
multitenancy
Posix shared filesystem support

Upvotes    Downvotes

8	0
7	0
7	0
7	0
6	0
6	0
3	0
2	0
2	0
1	0
2	1
1	0
0	0
1	1
0	1

15

Responses